

OpenCV Reference Manual

v2.3

December, 2011

Contents

<u>1 core. The Core Functionality</u>	<u>1</u>
<u>1.1 Basic Structures</u>	<u>1</u>
<u>cv::ocl::oclMat</u>	<u>1</u>
<u>cv::ocl::convertTo()</u>	<u>1</u>
<u>cv::ocl::copyTo()</u>	<u>2</u>
<u>cv::ocl::setTo()</u>	<u>3</u>
<u>1.2 operations on Arrays</u>	<u>4</u>
<u>cv::ocl::absdiff</u>	<u>4</u>
<u>cv::ocl::add</u>	<u>5</u>
<u>cv::ocl::bitwise_and</u>	<u>6</u>
<u>cv::ocl::bitwise_not</u>	<u>7</u>
<u>cv::ocl::bitwise_or</u>	<u>8</u>
<u>cv::ocl::bitwise_xor</u>	<u>9</u>
<u>cv::ocl::cartToPolar</u>	<u>10</u>
<u>cv::ocl::compare</u>	<u>11</u>
<u>cv::ocl::countNonZero</u>	<u>12</u>
<u>cv::ocl::divide</u>	<u>13</u>
<u>cv::ocl::exp</u>	<u>14</u>
<u>cv::ocl:: flip</u>	<u>15</u>
<u>cv::ocl::log</u>	<u>16</u>
<u>cv::ocl::LUT ()</u>	<u>17</u>
<u>cv::ocl::magnitude</u>	<u>18</u>
<u>cv::ocl::meanStdDev</u>	<u>19</u>
<u>cv::ocl::merge</u>	<u>20</u>
<u>cv::ocl:: minMax</u>	<u>21</u>
<u>cv::ocl:: minMaxLoc</u>	<u>22</u>
<u>cv::ocl::multiply</u>	<u>23</u>
<u>cv::ocl:: norm</u>	<u>24</u>
<u>cv::ocl::phase</u>	<u>25</u>
<u>cv::ocl::polarToCart</u>	<u>26</u>
<u>cv::ocl::split</u>	<u>27</u>
<u>cv::ocl::subtract</u>	<u>28</u>
<u>cv::ocl::transpose</u>	<u>29</u>
<u>2 imgproc. Image Processing</u>	<u>30</u>
<u>2.1 Histograms</u>	<u>30</u>
<u>cv::ocl::calcHist</u>	<u>30</u>
<u>2.2 Image Filtering</u>	<u>31</u>
<u>cv::ocl::bilateralFilter</u>	<u>31</u>
<u>cv::ocl::blur</u>	<u>32</u>
<u>cv::ocl::boxFilter</u>	<u>33</u>
<u>cv::ocl::copyMakeBorder</u>	<u>34</u>

[cv::ocl::dilate](#) 35
[cv::ocl::erode](#) 36
[cv::ocl::filter2D](#) 37
[cv::ocl::GaussianBlur](#) 39
[cv::ocl::Laplacian](#) 40
[cv::ocl::medianFilter \(\)\(pdf\)](#) 41
[cv::ocl::morphologyEx \(\)](#) 42
[cv::ocl::Scharr \(\)](#) 43
[cv::ocl::sepFilter2D](#) 44
[cv::ocl::Sobel](#) 45
[2.3 Geometric Image Transformations](#) 47
[cv::ocl::resize](#) 47
[cv::ocl::warpAffine](#) 49
[cv::ocl::warpPerspective](#) 50
[2.4 Miscellaneous Image Transformations](#) 51
[cv::ocl::cvtColor \(\)](#) 51
[cv::ocl::integral](#) 56
[cv::ocl::threshold](#) 58

1 1 core. The Core Functionality

1.1 1.1 Basic Structures

1.1.1 cv::ocl::oclMat

```

class oclMat
{
public:
    //default constructor
    oclMat();
    oclMat(int rows,int cols,int type);
    oclMat(Size size,int type);
    oclMat(int rows,int cols,int type,const Scalar & s);
    //copy constructor
    oclMat(const oclMat & m);
    //constructor for oclMatrix headers pointing to user-allocated data
    oclMat(int rows, int cols, int type, void *data, size_t step = Mat::AUTO_STEP);
    oclMat(Size size, int type, void *data, size_t step = Mat::AUTO_STEP);
    // creates a matrix header for a part of the bigger matrix
    oclMat(const oclMat &m, const Range &rowRange, const Range &colRange);
    oclMat(const oclMat &m, const Rect &roi);
    // builds oclMat from Mat. Perform blocking upload to device.
    explicit oclMat (const Mat &m);
    // destructor - calls release()
    oclMat();
    // assignment operators
  
```

```

oclMat &operator = (const oclMat &m);
// assignment operator. Perform blocking upload to device.
oclMat &operator = (const Mat &m);
// performs blocking upload data to oclMat.
void upload(const cv::Mat &m);
// downloads data from device to host memory. Blocking calls.
operator Mat() const;
void download(cv::Mat &m) const;
// returns a new oclMatrix header for the specified row
oclMat row(int y) const;
// returns a new oclMatrix header for the specified column
oclMat col(int x) const;
// ... for the specified row span
oclMat rowRange(int startrow, int endrow) const;
oclMat rowRange(const Range &r) const;
// ... for the specified column span
oclMat colRange(int startcol, int endcol) const;
oclMat colRange(const Range &r) const;
// returns deep copy of the oclMatrix, i.e. the data is copied
oclMat clone() const;
// copies the oclMatrix content to "m".
// It calls m.create(this->size(), this->type()).
// It supports any data type
void copyTo( oclMat &m ) const;
// copies those oclMatrix elements to "m" that are marked with non-zero
mask elements.
//It supports 8UC1 8UC4 32SC1 32SC4 32FC1 32FC4
void copyTo( oclMat &m, const oclMat &mask ) const;
// converts oclMatrix to another datatype with optional scaling. See cvCon-
vertScale.
//It supports 8UC1 8UC4 32SC1 32SC4 32FC1 32FC4
void convertTo( oclMat &m, int rtype, double alpha = 1, double beta = 0 )
const;
void assignTo( oclMat &m, int type = -1 ) const;
// sets every oclMatrix element to s
//It supports 8UC1 8UC4 32SC1 32SC4 32FC1 32FC4
oclMat &operator = (const Scalar &s);
// sets some of the oclMatrix elements to s, according to the mask
//It supports 8UC1 8UC4 32SC1 32SC4 32FC1 32FC4
oclMat &setTo(const Scalar &s, const oclMat &mask = oclMat());
// creates alternative oclMatrix header for the same data, with different
// number of channels and/or different number of rows. see cvReshape.
oclMat reshape(int cn, int rows = 0) const;
// allocates new oclMatrix data unless the oclMatrix already has specified
size and type.
// previous data is unreferenced if needed.

```

```

void create(int rows, int cols, int type);
void create(Size size, int type);
// decreases reference counter;
// deallocate the data when reference counter reaches 0.
void release();
// swaps with other smart pointer
void swap(oclMat &mat);
// locates oclMatrix header within a parent oclMatrix. See below
void locateROI( Size &wholeSize, Point &ofs ) const;
// moves/resizes the current oclMatrix ROI inside the parent oclMatrix.
oclMat &adjustROI( int dtop, int dbottom, int dleft, int dright );
// extracts a rectangular sub-oclMatrix
// (this is a generalized form of row, rowRange etc.)
oclMat operator()( Range rowRange, Range colRange ) const;
oclMat operator()( const Rect &roi ) const;
// returns true if the oclMatrix data is continuous
// (i.e. when there are no gaps between successive rows).
// similar to CV_IS_oclMat_CONT(cvoclMat->type)
bool isContinuous() const;
// returns element size in bytes,
// similar to CV_ELEM_SIZE(cvMat->type)
size_t elemSize() const;
// returns the size of element channel in bytes.
size_t elemSize1() const;
// returns element type, similar to CV_MAT_TYPE(cvMat->type)
int type() const;
// returns element type, similar to CV_MAT_DEPTH(cvMat->type)
int depth() const;
// returns element type, similar to CV_MAT_CN(cvMat->type)
int channels() const;
// returns step/elemSize1()
size_t step1() const;
// returns oclMatrix size:
// width == number of columns, height == number of rows
Size size() const;
// returns true if oclMatrix data is NULL
bool empty() const;
// returns pointer to y-th row
uchar *ptr(int y = 0);
const uchar *ptr(int y = 0) const;
// matrix transposition
oclMat t() const;
/*! includes several bit-fields:
- the magic signature
- continuity flag
- depth

```

```

- number of channels
*/
int flags;
// the number of rows and columns
int rows, cols;
// a distance between successive rows in bytes; includes the gap if any
size_t step;
// pointer to the data(OCL memory object)
uchar *data;
// pointer to the reference counter;
// when oclMatrix points to user-allocated data, the pointer is NULL
int *refcount;
// helper fields used in locateROI and adjustROI
//datastart and dataend are not used in current version
uchar *datastart;
uchar *dataend;
// OpenCL context associated with the oclMat object.
Context *clCxt;
//add offset for handle ROI, calculated in byte
int offset;
//add wholerows and wholecols for the whole matrix, datastart and dataend
are no longer used
int wholerows;
int wholecols;
//add download_channels for 3 channels to 4 channels
int download_channels;
}

```

1.1.2 cv::ocl::convertTo()

Converts array to another datatype with optional scaling.

Supports ().

m The destination matrix. If it does not have a proper size or type before the operation, it will be reallocated

rtype The desired destination matrix type, or rather, the depth (since the number of channels will be the same with the source one). If rtype is negative, the destination matrix will have the same type as the source.

alpha The optional scale factor

beta The optional delta, added to the scaled values.

The method converts source pixel values to the target datatype. saturate cast<> is applied in the end to avoid possible overflows:

$$m(x, y) = \text{saturate_cast} < rType > (\alpha(*this)(x, y) + \beta)$$

1.1.3 `cv::ocl::copyTo()`

Copies the matrix to another one.

`Supports()`.

m The destination matrix. If it does not have a proper size or type before the operation, it will be reallocated

mask The operation mask. Its non-zero elements indicate, which matrix elements need to be copied

The method copies the matrix data to another matrix. Before copying the data, the method invokes

so that the destination matrix is reallocated if needed. While `m.copyTo(m);` will work as expected, i.e. will have no effect, the function does not handle the case of a partial overlap between the source and the destination matrices.

When the operation mask is specified, and the `Mat::create` call shown above reallocated the matrix, the newly allocated matrix is initialized with all 0's before copying the data.

1.1.4 `cv::ocl::setTo()`

Sets all or some of the array elements to the specified value.

Supports().

s Assigned scalar, which is converted to the actual array type

mask The operation mask of the same size as `*this`

This is the advanced variant of `Mat::operator=(const Scalar& s)` operator.

1.2 1.2 operations on Arrays

1.2.1 cv::ocl::absdiff

Computes per-element absolute difference between two arrays or between array and a scalar.

Supports all data types except CV_8SC1, CV_8SC2, CV_8SC3 and CV_8SC4.

src1 The first input array

src2 The second input array, must be the same size and same type as src1

sc Scalar, the second input parameter

dst The destination array, it will have the same size and same type as src1

The functions absdiff compute:

- absolute difference between two arrays

$$\text{dst}(\text{I}) = \text{saturate}(|\text{src1}(\text{I}) - \text{src2}(\text{I})|)$$

- or absolute difference between array and a scalar:

$$\text{dst}(\text{I}) = \text{saturate}(|\text{src1}(\text{I}) - \text{sc}|)$$

Where I is multi-dimensional index of array elements. In the case of multi-channel arrays each channel is processed independently.

1.2.2 cv::ocl::add

Computes the per-element sum of two array and a scalar.

Supports all data types except CV_8SC1, CV_8SC2, CV_8SC3 and CV_8SC4.

src1 The first source array

src2 The second source array. It must have the same size and same type as src1

sc Scalar; the second input parameter

dst The destination array; it will have the same size and same type as src1

mask The optional operation mask, 8-bit single channel array; specifies elements of the destination array to be changed

The functions add compute:

- the sum of two arrays:

$$\text{dst}(\text{I}) = \text{saturate}(\text{src1}(\text{I}) + \text{src2}(\text{I})) \text{ if } \text{mask}(\text{I}) \neq 0$$

- or the sum of array and a scalar:

$$\text{dst}(\text{I}) = \text{saturate}(\text{src1}(\text{I}) + \text{sc}) \text{ if } \text{mask}(\text{I}) \neq 0$$

where I is multi-dimensional index of array elements.

The first function in the above list can be replaced with matrix expressions: in the case of multi-channel arrays each channel is processed independently.

1.2.3 cv::ocl::bitwise_and

Calculates per-element bit-wise conjunction of two arrays and an array and a scalar.

Supports all data types except CV_64FC1, CV_64FC2, CV_64FC3 and CV_64FC4.

src1 The first source array

src2 The second source array. It must have the same size and same type as src1

sc Scalar; the second input parameter

dst The destination array; it will have the same size and same type as src1

mask The optional operation mask, 8-bit single channel array; specifies elements of the destination array to be changed

The functions bitwise and compute per-element bit-wise logical conjunction:

- of two arrays:

$$\text{dst}(I) = \text{src1}(I) \wedge \text{src2}(I) \text{ if } \text{mask}(I) \neq 0$$

- or array and a scalar:

$$\text{dst}(I) = \text{src1}(I) \wedge \text{sc} \text{ if } \text{mask}(I) \neq 0$$

In the case of floating-point arrays their machine-specific bit representations (usually IEEE754-compliant) are used for the operation, and in the case of multi-channel arrays each channel is processed independently.

1.2.4 cv::ocl::bitwise_not

Inverts every bit of array.

Supports CV_8U, CV_8S, CV_16U, CV_16S, CV_32S, CV_32F data types.

src1 The source array

dst The destination array; it is reallocated to be of the same size and the same type as src

mask The optional operation mask, 8-bit single channel array; specifies elements of the destination array to be changed

The functions bitwise not compute per-element bit-wise inversion of the source array:

$$\text{dst}(I) = \neg \text{src}(I)$$

In the case of floating-point source array its machine-specific bit representation (usually IEEE754-compliant) is used for the operation. in the case of multi-channel arrays each channel is processed independently.

1.2.5 cv::ocl::bitwise_or

Calculates per-element bit-wise disjunction of two arrays and an array and a scalar.

Supports all data types except CV_64FC1, CV_64FC2, CV_64FC3 and CV_64FC4.

src1 The first source array

src2 The second source array. It must have the same size and same type as src1

sc Scalar; the second input parameter

dst The destination array; it is reallocated to be of the same size and the same type as src1

mask The optional operation mask, 8-bit single channel array; specifies elements of the destination array to be changed

The functions bitwise_or compute per-element bit-wise logical disjunction

- of two arrays:

$$\text{dst}(I) = \text{src1}(I) \vee \text{src2}(I) \text{ if } \text{mask}(I) \neq 0$$

- or array and a scalar:

$$\text{dst}(I) = \text{src1}(I) \vee \text{sc} \text{ if } \text{mask}(I) \neq 0$$

In the case of floating-point arrays their machine-specific bit representations (usually IEEE754-compliant) are used for the operation. in the case of multi-channel arrays each channel is processed independently.

1.2.6 cv::ocl::bitwise_xor

Calculates per-element bit-wise “exclusive or” operation on two arrays and an array and a scalar.

Supports all data types except CV_64FC1,CV_64FC2,CV_64FC3 and CV_64FC4.

src1 The first source array

src2 The second source array. It must have the same size and same type as src1

sc Scalar; the second input parameter

dst The destination array; it is reallocated to be of the same size and the same type as src1

mask The optional operation mask, 8-bit single channel array; specifies elements of the destination array to be changed

The functions bitwise xor compute per-element bit-wise logical ”exclusive or” operation

- on two arrays:

$$\text{dst}(I) = \text{src1}(I) \oplus \text{src2}(I) \text{ if } \text{mask}(I) \neq 0$$

- or array and a scalar:

$$\text{dst}(I) = \text{src1}(I) \oplus \text{sc if } \text{mask}(I) \neq 0$$

In the case of floating-point arrays their machine-specific bit representations (usually IEEE754-compliant) are used for the operation. in the case of multi-channel arrays each channel is processed independently.

1.2.7 cv::ocl::cartToPolar

Calculates the magnitude and angle of 2d vectors.

Supports only CV_32F, CV_64F data types.

x The array of x-coordinates; must be single-precision or double-precision floating-point array

y The array of y-coordinates; it must have the same size and same type as x

magnitude The destination array of magnitudes of the same size and same type as x

angle The destination array of angles of the same size and same type as x. The angles are measured in radians (0 to 2π) or in degrees (0 to 360 degrees).

angleInDegrees The flag indicating whether the angles are measured in radians, which is default mode, or in degrees

The function `cartToPolar` calculates either the magnitude, angle, or both of every 2d vector (`x(I),y(I)`):

$$\text{magnitude}(I) = \sqrt{x(I)^2 + y(I)^2}$$

$$\text{angle}(I) = \text{atan2}(y(I), x(I)) \cdot [180/\pi]$$

The angles are calculated with $\sim 0.3^\circ$ accuracy. For the (0,0) point, the angle is set to 0.

1.2.8 cv::ocl::compare

Performs per-element comparison of two arrays or an array and scalar value.

Supports all data types except CV_8SC1,CV_8SC2,CV_8SC3,CV_8SC4.

src1 The first source array

src2 The second source array; must have the same size and same type as src1

dst The destination array; will have the same size as src1 and type=CV_8UC1

cmpop The flag specifying the relation between the elements to be checked

CMP_EQ src1(I) = src2(I) or src1(I) = value

CMP_GT src1(I) > src2(I) or src1(I) > value

CMP_GE src1(I) ≥ src2(I) or src1(I) ≥ value

CMP_LT src1(I) < src2(I) or src1(I) < value

CMP_LE src1(I) ≤ src2(I) or src1(I) ≤ value

CMP_NE src1(I) ≠ src2(I) or src1(I) ≠ value

The functions compare each element of src1 with the corresponding element of src2 or with real scalar value. When the comparison result is true, the corresponding element of destination array is set to 255, otherwise it is set to 0:

- $\text{dst}(\text{I}) = \text{src1}(\text{I}) \text{ cmpop } \text{src2}(\text{I}) ? 255 : 0$
- $\text{dst}(\text{I}) = \text{src1}(\text{I}) \text{ cmpop value} ? 255 : 0$

The comparison operations can be replaced with the equivalent matrix expressions:

1.2.9 cv::ocl::countNonZero

Counts non-zero array elements.

Supports all data types.

mtx Single-channel array

The function cvCountNonZero returns the number of non-zero elements in
mtx:

$$\sum_{I:\text{mtx}(I)\neq 0} 1$$

1.2.10 cv::ocl::divide

Performs per-element division of two arrays or a scalar by an array.

Supports all data types except CV_8SC1, CV_8SC2, CV_8SC3 and CV_8SC4.

src1 The first source array

src2 The second source array; should have the same size and same type as src1

scale Scale factor

dst The destination array; will have the same size and same type as src2

The functions divide compute:

- divide one array by another:

$$\text{dst}(I) = \text{saturate}(\text{src1}(I) * \text{scale} / \text{src2}(I))$$

- or a scalar by array, when there is no src1:

$$\text{dst}(I) = \text{saturate}(\text{scale} / \text{src2}(I))$$

The result will have the same type as src1. When src2(I)=0, dst(I)=0 too.

1.2.11 `cv::ocl::exp`

Calculates the exponent of every array element.

Supports only CV_32FC1 data type.

src The source array

dst The destination array; will have the same size and same type as src The function `exp` calculates the exponent of every element of the input array:

$$\text{dst}[I] = e^{\text{src}}(I)$$

The maximum relative error is about 7×10^{-6} for single-precision and less than 10^{-10} for doubleprecision. Currently, the function converts denormalized values to zeros on output. Special values (NaN, $\pm\infty$) are not handled.

1.2.12 cv::ocl:: flip

Flips a 2D array around vertical, horizontal or both axes.

Supports all data types.

src The source array

dst The destination array; will have the same size and same type as src

flipCode Specifies how to flip the array: 0 means flipping around the x-axis, positive (e.g., 1) means flipping around y-axis, and negative (e.g., -1) means flipping around both axes. See also the discussion below for the formulas.

The function flip flips the array in one of three different ways (row and column indices are 0-based):

$$dst_{ij} = \begin{cases} src_{src.rows-i-1, j} & \text{if flipCode} = 0 \\ src_{src.i, src.cols-j-1} & \text{if flipCode} > 0 \\ src_{src.rows-i-1, src.cols-j-1} & \text{if flipCode} < 0 \end{cases}$$

The example scenarios of function use are:

- vertical flipping of the image (flipCode = 0) to switch between top-left and bottom-left image origin, which is a typical operation in video processing in Windows.
- horizontal flipping of the image with subsequent horizontal shift and absolute difference calculation to check for a vertical-axis symmetry (flipCode > 0)
- simultaneous horizontal and vertical flipping of the image with subsequent shift and absolute difference calculation to check for a central symmetry (flipCode < 0)
- reversing the order of 1d point arrays (flipCode > 0 or flipCode = 0)

1.2.13 `cv::ocl::log`

Calculates the natural logarithm of every array element.

Supports only CV_32FC1 data type.

src The source array

dst The destination array; will have the same size and same type as src

The function `log` calculates the natural logarithm of the absolute value of every element of the input array:

$$\text{dst}(I) = \begin{cases} \log |\text{src}(I)| & \text{if } \text{src}(I) \neq 0 \\ C & \text{otherwise} \end{cases}$$

Where C is a large negative number (about -700 in the current implementation). The maximum relative error is about 7×10^{-6} for single-precision input and less than 10^{-10} for double-precision input. Special values (NaN, $\pm\infty$) are not handled.

1.2.14 `cv::ocl::LUT` ()

Performs a look-up table transform of an array.

Supports()

src Source array of 8-bit elements

lut Look-up table of 256 elements. In the case of multi-channel source array, the table should either have a single channel (in this case the same table is used for all channels) or the same number of channels as in the source array

dst Destination array; will have the same size and the same number of channels as src, and the same depth as lut

The function LUT fills the destination array with values from the look-up table. Indices of the entries are taken from the source array. That is, the function processes each element of src as follows:

$$\text{dst}(I) \leftarrow \text{lut}(\text{src}(I) + d)$$

where

$$d = \begin{cases} 0 & \text{if src has depth CV_8U} \\ 128 & \text{if src has depth CV_8S} \end{cases}$$

1.2.15 `cv::ocl::magnitude`

Calculates magnitude of 2D vectors.

Supports only CV_32F, CV_64F data types.

x The floating-point array of x-coordinates of the vectors

y The floating-point array of y-coordinates of the vectors; must have the same size as x

magnitude The destination array; will have the same size and same type as x

The function `magnitude` calculates magnitude of 2D vectors formed from the corresponding elements of x and y arrays:

$$\text{dst}(I) = \sqrt{x(I)^2 + y(I)^2}$$

1.2.16 cv::ocl::meanStdDev

Calculates mean and standard deviation of array elements.

Supports all data types except CV_32F,CV_64F.

mtx The source array; it should have 1 to 4 channels

mean The output parameter: computed mean value

stddev The output parameter: computed standard deviation

The functions meanStdDev compute the mean and the standard deviation M of array elements, independently for each channel, and return it via the output parameters:

$$\begin{aligned} N &= \sum_I 1 \\ mean_c &= \frac{\sum_I src(I)_c}{N} \\ stddev_c &= \sqrt{\sum_I (src(I)_c - mean_c)^2} \end{aligned}$$

1.2.17 `cv::ocl::merge`

Composes a multi-channel array from several single-channel arrays.

Supports all data types.

mv The source array or vector of the single-channel matrices to be merged.

All the matrices in mv must have the same size and the same type

count The number of source matrices when mv is a plain C array; must be greater than zero

dst The destination array; will have the same size and the same depth as mv[0], the number of channels will match the number of source matrices

The functions merge merge several single-channel arrays (or rather interleave their elements) to make a single multi-channel array.

$$\text{dst}(I)_c = \text{mv}[c](I)$$

The function `cv::split` does the reverse operation and if you need to merge several multichannel images or shuffle channels in some other advanced way.

1.2.18 cv::ocl:: minMax

Calculates per-element minimum and maximum of two arrays or array and a scalar.

Supports all data types.

Finds global minimum and maximum in a whole array or sub-array

src The source single-channel array

minVal Pointer to returned minimum value; NULL if not required

maxVal Pointer to returned maximum value; NULL if not required

mask The optional mask used to select a sub-array

The functions minMax find minimum and maximum element values. The extremums are searched across the whole array, or, if mask is not an empty array, in the specified array region.

The functions do not work with multi-channel arrays.

In the case of a sparse matrix the minimum is found among non-zero elements only.

1.2.19 cv::ocl:: minMaxLoc

Finds global minimum and maximum in a whole array or sub-array

Supports all data types.

src The source single-channel array

minVal Pointer to returned minimum value; NULL if not required

maxVal Pointer to returned maximum value; NULL if not required

minLoc Pointer to returned minimum location (in 2D case); NULL if not required

maxLoc Pointer to returned maximum location (in 2D case); NULL if not required

mask The optional mask used to select a sub-array

The functions minMaxLoc find minimum and maximum element values and their positions. The extremums are searched across the whole array, or, if mask is not an empty array, in the specified array region.

The functions do not work with multi-channel arrays.

In the case of a sparse matrix the minimum is found among non-zero elements only.

1.2.20 `cv::ocl::multiply`

Calculates the per-element scaled product of two arrays.

Supports all data types except CV_8SC1, CV_8SC2, CV_8SC3 and CV_8SC4.

src1 The first source array

src2 The second source array of the same size and the same type as src1

dst The destination array; will have the same size and the same type as src1

scale The optional scale factor

The function multiply calculates the per-element product of two arrays:

$\text{dst}(I) = \text{saturate}(\text{scale} \cdot \text{src1}(I) \cdot \text{src2}(I))$

There is also Matrix Expressions -friendly variant of the first function.

1.2.21 cv::ocl:: norm

Calculates absolute array norm, absolute difference norm, or relative difference norm.

Supports only CV_8UC1 data type.

src1 The first source array

src2 The second source array of the same size and the same type as src1

normType Type of the norm; see the discussion below

The functions norm calculate the absolute norm of src1 (when there is no src2):

$$norm = \begin{cases} \|src1\|_{L_\infty} = \max_I |src1(I)| & \text{if normType} = \text{NORM_INF} \\ \|src1\|_{L_1} = \sum_I |src1(I)| & \text{if normType} = \text{NORM_L1} \\ \|src1\|_{L_2} = \sqrt{\sum_I src1(I)^2} & \text{if normType} = \text{NORM_L2} \end{cases}$$

or an absolute or relative difference norm if src2 is there:

$$norm = \begin{cases} \|src1 - src2\|_{L_\infty} = \max_I |src1(I) - src2(I)| & \text{if normType} = \text{NORM_INF} \\ \|src1 - src2\|_{L_1} = \sum_I |src1(I) - src2(I)| & \text{if normType} = \text{NORM_L1} \\ \|src1 - src2\|_{L_2} = \sqrt{\sum_I (src1(I) - src2(I))^2} & \text{if normType} = \text{NORM_L2} \end{cases}$$

or

$$norm = \begin{cases} \frac{\|src1 - src2\|_{L_\infty}}{\|src2\|_{L_\infty}} & \text{if normType} = \text{NORM_INF} \\ \frac{\|src1 - src2\|_{L_1}}{\|src2\|_{L_1}} & \text{if normType} = \text{NORM_L1} \\ \frac{\|src1 - src2\|_{L_2}}{\|src2\|_{L_2}} & \text{if normType} = \text{NORM_L2} \end{cases}$$

The functions norm return the calculated norm.

When there is mask parameter, and it is not empty (then it should have type CV_8U and the same size as src1), the norm is computed only over the specified by the mask region.

A multiple-channel source arrays are treated as a single-channel, that is, the results for all channels are combined.

1.2.22 cv::ocl::phase

Calculates the rotation angle of 2d vectors.

Supports only CV_32F CV_64F data types.

x The source floating-point array of x-coordinates of 2D vectors

y The source array of y-coordinates of 2D vectors; must have the same size and the same type as x

angle The destination array of vector angles; it will have the same size and same type as x

angleInDegrees When it is true, the function will compute angle in degrees, otherwise they will be measured in radians

The function phase computes the rotation angle of each 2D vector that is formed from the corresponding elements of x and y:

$$\text{angle}(I) = \text{atan2}(y(I), x(I))$$

The angle estimation accuracy is 0.3° , when $x(I)=y(I)=0$, the corresponding $\text{angle}(I)$ is set to 0.

1.2.23 cv::ocl::polarToCart

Computes x and y coordinates of 2D vectors from their magnitude and angle.

Supports only CV_32F CV_64F data type.

magnitude The source floating-point array of magnitudes of 2D vectors. It can be an empty matrix (=Mat()) - in this case the function assumes that all the magnitudes are =1. If it's not empty, it must have the same size and same type as angle

angle The source floating-point array of angles of the 2D vectors

x The destination array of x-coordinates of 2D vectors; will have the same size and the same type as angle

y The destination array of y-coordinates of 2D vectors; will have the same size and the same type as angle

angleInDegrees When it is true, the input angles are measured in degrees, otherwise they are measured in radians

The function polarToCart computes the cartesian coordinates of each 2D vector represented by the corresponding elements of magnitude and angle:

$$\begin{aligned}x(I) &= \text{magnitude}(I) \cos(\text{angle}(I)) \\y(I) &= \text{magnitude}(I) \sin(\text{angle}(I))\end{aligned}$$

The relative accuracy of the estimated coordinates is 10^{-6} .

1.2.24 `cv::ocl::split`

Divides multi-channel array into several single-channel arrays.

Supports all data types.

mtx The source multi-channel array

mv The destination array or vector of arrays; The number of arrays must match `mtx.channels()`. The arrays themselves will be reallocated if needed

The functions `split` split multi-channel array into separate single-channel arrays:

$$\text{mv}[c](I) = \text{mtx}(I)_c$$

1.2.25 cv::ocl::subtract

Calculates per-element difference between two arrays or array and a scalar.

Supports all data types except CV_8SC1, CV_8SC2, CV_8SC3 and CV_8SC4.

src1 The first source array

src2 The second source array. It must have the same size and same type as src1

sc Scalar; the first or the second input parameter

dst The destination array; it will have the same size and same type as src1

mask The optional operation mask, 8-bit single channel array; specifies elements of the destination array to be changed

The functions subtract compute:

- the difference between two arrays:

$$\text{dst}(I) = \text{saturate}(\text{src1}(I) - \text{src2}(I)) \text{ if } \text{mask}(I) \neq 0$$

- the difference between array and a scalar:

$$\text{dst}(I) = \text{saturate}(\text{src1}(I) - \text{sc}) \text{ if } \text{mask}(I) \neq 0$$

- the difference between scalar and an array:

$$\text{dst}(I) = \text{saturate}(\text{sc} - \text{src2}(I)) \text{ if } \text{mask}(I) \neq 0$$

where I is multi-dimensional index of array elements.

The first function in the above list can be replaced with matrix expressions:

1.2.26 `cv::ocl::transpose`

Transposes a matrix.

Supports CV_8UC1, 8UC4, 8SC4, 16UC2, 16SC2, 32SC1 and 32FC1 data types.

src The source array

dst The destination array of the same type as src

The function `cv::transpose` transposes the matrix src:

$\text{dst}(i, j) = \text{src}(j, i)$

Note that no complex conjugation is done in the case of a complex matrix, it should be done separately if needed.

2 2 imgproc. Image Processing

2.1 2.1 Histograms

2.1.1 cv::ocl::calcHist

Calculates histogram of one or more arrays

Supports only 8UC1 data type.

arrays Source arrays. They all should have the same depth, CV 8U, and the same size. Each of them can have an arbitrary number of channels.

hist The output histogram, a dense or sparse dims-dimensional array.

2.2 2.2 Image Filtering

2.2.1 cv::ocl::bilateralFilter

Applies bilateral filter to the image.

Supports 8UC1 8UC4 data types.

src The source 8-bit or floating-point, 1-channel or 3-channel image

dst The destination image; will have the same size and the same type as src

d The diameter of each pixel neighborhood, that is used during filtering. If it is non-positive, it's computed from sigmaSpace

sigmaColor Filter sigma in the color space. Larger value of the parameter means that farther colors within the pixel neighborhood (see sigmaSpace) will be mixed together, resulting in larger areas of semi-equal color

sigmaSpace Filter sigma in the coordinate space. Larger value of the parameter means that farther pixels will influence each other (as long as their colors are close enough; see sigmaColor). Then $d > 0$, it specifies the neighborhood size regardless of sigmaSpace, otherwise d is proportional to sigmaSpace.

2.2.2 cv::ocl::blur

Smooths image using normalized box filter

Supports data type: CV_8UC1, CV_8UC4, CV_32FC1 and CV_32FC4

src The source image

dst The destination image; will have the same size and the same type as src

ksize The smoothing kernel size

anchor The anchor point. The default value Point(-1,-1) means that the anchor is at the kernel center

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

The function smooths the image using the kernel:

$$K = \frac{1}{\text{ksize.width} * \text{ksize.height}} \begin{bmatrix} 1 & 1 & 1 & \dots & 1 & 1 \\ 1 & 1 & 1 & \dots & 1 & 1 \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & 1 & 1 & \dots & 1 & 1 \end{bmatrix}$$

The call blur(src, dst, ksize, anchor, borderType) is equivalent to boxFilter(src, dst, src.type(), anchor, true, borderType).

2.2.3 cv::ocl::boxFilter

Smooths image using box filter

Supports data type: CV_8UC1, CV_8UC4, CV_32FC1 and CV_32FC4.

src The source image

dst The destination image; will have the same size and the same type as src

ksize The smoothing kernel size

anchor The anchor point. The default value Point(-1,-1) means that the anchor is at the kernel center

normalize Indicates, whether the kernel is normalized by its area or not

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

The function smooths the image using the kernel:

$$K = \alpha \begin{bmatrix} 1 & 1 & 1 & \dots & 1 & 1 \\ 1 & 1 & 1 & \dots & 1 & 1 \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & 1 & 1 & \dots & 1 & 1 \end{bmatrix}$$

where

$$\alpha = \begin{cases} \frac{1}{\text{ksize.width} * \text{ksize.height}} & \text{when normalize=true} \\ 1 & \text{otherwise} \end{cases}$$

Unnormalized box filter is useful for computing various integral characteristics over each pixel neighborhood, such as covariation matrices of image derivatives (used in dense optical flow algorithms etc.). If you need to compute pixel sums over variable-size windows, use cv::integral.

2.2.4 cv::ocl::copyMakeBorder

Forms a border around the image.

Supports CV_8UC1, CV_8UC4, CV_32SC1 data types.

src The source image

dst The destination image; will have the same type as src and the size `size(src.cols+left+right, src.rows+top+bottom)`

top, bottom, left, right Specify how much pixels in each direction from the source image rectangle one needs to extrapolate, e.g. `top=1, bottom=1, left=1, right=1` mean that 1 ixel-wide border needs to be built

borderType The border type

value The border value if `borderType==BORDER_CONSTANT`

The function copies the source image into the middle of the destination image. The areas to the left, to the right, above and below the copied source image will be filled with extrapolated pixels. This is not what `cv::FilterEngine` or based on it filtering functions do (they extrapolate pixels on-fly), but what other more complex functions, including your own, may do to simplify image boundary handling.

The function supports the mode when src is already in the middle of dst. In this case the function does not copy src itself, but simply constructs the border, e.g.:

2.2.5 cv::ocl::dilate

Dilates an image by using a specific structuring element.

Supports data types: CV_8UC1, CV_8UC4, CV_32FC1 and CV_32FC4.

src The source image

dst The destination image. It will have the same size and the same type as src

element The structuring element used for dilation. If element=Mat(), a 3×3 rectangular structuring element is used

anchor Position of the anchor within the element. The default value (-1, -1) means that the anchor is at the element center

iterations The number of times dilation is applied

The function dilates the source image using the specified structuring element that determines the shape of a pixel neighborhood over which the maximum is taken:

$$\text{dst}(x, y) = \max_{(x', y') : \text{element}(x', y') \neq 0} \text{src}(x + x', y + y')$$

The function supports the in-place mode. Dilation can be applied several (iterations) times. In the case of multi-channel images each channel is processed independently.

2.2.6 cv::ocl::erode

Erodes an image by using a specific structuring element.

Supports data types: CV_8UC1, CV_8UC4, CV_32FC1 and CV_32FC4.

src The source image

dst The destination image. It will have the same size and the same type as src

element The structuring element used for dilation. If element=Mat(), a 3×3 rectangular structuring element is used

anchor Position of the anchor within the element. The default value (-1, -1) means that the anchor is at the element center

iterations The number of times erosion is applied

The function erodes the source image using the specified structuring element that determines the shape of a pixel neighborhood over which the minimum is taken:

$$\text{dst}(x, y) = \min_{(x', y') : \text{element}(x', y') \neq 0} \text{src}(x + x', y + y')$$

The function supports the in-place mode. Erosion can be applied several (iterations) times. In the case of multi-channel images each channel is processed independently.

2.2.7 cv::ocl::filter2D

Convolves an image with the kernel

Supports

src The source image

dst The destination image. It will have the same size and the same number of channels as src

ddepth The desired depth of the destination image. If it is negative, it will be the same as src.depth()

kernel Convolution kernel (or rather a correlation kernel), a single-channel floating point matrix. If you want to apply different kernels to different channels, split the image into separate color planes using cv::split and process them individually

anchor The anchor of the kernel that indicates the relative position of a filtered point within the kernel. The anchor should lie within the kernel. The special default value (-1,-1) means that the anchor is at the kernel center

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

The function applies an arbitrary linear filter to the image. In-place operation is supported. When the aperture is partially outside the image, the function interpolates outlier pixel values according to the specified border mode.

The function does actually computes correlation, not the convolution:

$$\text{dst}(x, y) = \sum_{\substack{0 \leq x' < \text{kernel.cols} \\ 0 \leq y' < \text{kernel.rows}}} \text{kernel}(x', y') * \text{src}(x + x' - \text{anchor.x}, y + y' - \text{anchor.y})$$

That is, the kernel is not mirrored around the anchor point. If you need a real convolution, flip the kernel using cv::flip and set the new anchor to (kernel.cols - anchor.x - 1, kernel.rows - anchor.y - 1).

The function uses DFT-based algorithm in case of sufficiently large kernels (11×11) and the direct algorithm for small kernels.

2.2.8 cv::ocl::GaussianBlur

Smooths image using a Gaussian filter

Supports

src The source image

dst The destination image; will have the same size and the same type as src

ksize The Gaussian kernel size; ksize.width and ksize.height can differ, but they both must be positive and odd. Or, they can be zero's, then they are computed from sigma*

sigmaX, sigmaY The Gaussian kernel standard deviations in X and Y direction. If sigmaY is zero, it is set to be equal to sigmaX. If they are both zeros, they are computed from ksize.width and ksize.height. To fully control the result regardless of possible future modification of all this semantics, it is recommended to specify all of ksize, sigmaX and sigmaY

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

The function convolves the source image with the specified Gaussian kernel. In-place filtering is supported.

2.2.9 cv::ocl::Laplacian

Calculates the Laplacian of an image

Supports only ksize = 1 and ksize = 3 8UC1 8UC4 32FC1 32FC4 data types.

src Source image

dst Destination image; will have the same size and the same number of channels as src

ddepth The desired depth of the destination image

ksize The aperture size used to compute the second-derivative filters. It must be positive and odd

scale The optional scale factor for the computed Laplacian values (by default, no scaling is applied)

The function calculates the Laplacian of the source image by adding up the second x and y derivatives calculated using the Sobel operator:

$$\text{dst} = \Delta_{\text{src}} = \frac{\partial^2 \text{src}}{\partial x^2} + \frac{\partial^2 \text{src}}{\partial y^2}$$

This is done when ksize > 1. When ksize == 1, the Laplacian is computed by filtering the image with the following 3×3 aperture:

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

2.2.10 `cv::ocl::medianFilter()`(pdf)

Smooths image using median filter

Supports ().

src The source 1-, 3- or 4-channel image. When ksize is 3 or 5, the image depth should be CV_8U, CV_16U or CV_32F. For larger aperture sizes it can only be CV_8U

dst The destination array; will have the same size and the same type as src

ksize The aperture linear size. It must be odd and more than 1, i.e. 3, 5, 7

...

The function smooths image using the median filter with $\text{ksize} \times \text{ksize}$ aperture. Each channel of a multi-channel image is processed independently. In-place operation is supported.

2.2.11 cv::ocl::morphologyEx ()

Performs advanced morphological transformations.

Supports()

src Source image

dst Destination image. It will have the same size and the same type as src

element Structuring element

op Type of morphological operation, one of the following:

MORPH_OPEN opening

MORPH_CLOSE closing

MORPH_GRADIENT morphological gradient

MORPH_TOPHAT “top hat”

MORPH_BLACKHAT “black hat”

iterations Number of times erosion and dilation are applied

The function can perform advanced morphological transformations using erosion and dilation as basic operations.

Opening:

$$\text{dst} = \text{open}(\text{src}, \text{element}) = \text{dilate}(\text{erode}(\text{src}, \text{element}))$$

Closing:

$$\text{dst} = \text{close}(\text{src}, \text{element}) = \text{erode}(\text{dilate}(\text{src}, \text{element}))$$

Morphological gradient:

$$\text{dst} = \text{morph_grad}(\text{src}, \text{element}) = \text{dilate}(\text{src}, \text{element}) - \text{erode}(\text{src}, \text{element})$$

“Top hat”:

$$\text{dst} = \text{tophat}(\text{src}, \text{element}) = \text{src} - \text{open}(\text{src}, \text{element})$$

“Black hat”:

$$\text{dst} = \text{blackhat}(\text{src}, \text{element}) = \text{close}(\text{src}, \text{element}) - \text{src}$$

Any of the operations can be done in-place.

2.2.12 cv::ocl::Scharr ()

Calculates the first x- or y- image derivative using Scharr operator

Supports().

src The source image

dst The destination image; will have the same size and the same number of channels as src

ddepth The destination image depth

xorder Order of the derivative x

yorder Order of the derivative y

scale The optional scale factor for the computed derivative values (by default, no scaling is applied)

delta The optional delta value, added to the results prior to storing them in dst

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

The function computes the first x- or y- spatial image derivative using Scharr operator. The call

```
Scharr(src, dst, ddepth, xorder, yorder, scale, delta, borderType).
```

is equivalent to

```
Sobel(src, dst, ddepth, xorder, yorder, CV_SCHARR, scale, delta,  
borderType).
```

2.2.13 cv::ocl::sepFilter2D

Applies separable linear filter to an image

Supports .

src The source image

dst The destination image; will have the same size and the same number of channels as src

ddepth The destination image depth

rowKernel The coefficients for filtering each row

columnKernel The coefficients for filtering each column

anchor The anchor position within the kernel; The default value (-1, 1) means that the anchor is at the kernel center

delta The value added to the filtered results before storing them

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

The function applies a separable linear filter to the image. That is, first, every row of src is filtered with 1D kernel rowKernel. Then, every column of the result is filtered with 1D kernel columnKernel and the final result shifted by delta is stored in dst.

2.2.14 cv::ocl::Sobel

Calculates the first, second, third or mixed image derivatives using an extended Sobel operator

Supports .

src The source image

dst The destination image; will have the same size and the same number of channels as src

ddepth The destination image depth

xorder Order of the derivative x

yorder Order of the derivative y

ksize Size of the extended Sobel kernel, must be 1, 3, 5 or 7

scale The optional scale factor for the computed derivative values (by default, no scaling is applied)

delta The optional delta value, added to the results prior to storing them in dst

borderType The border mode used to extrapolate pixels outside of the image, the details are shown as follows:

- BORDER_REPLICATE: aaaaaa|abcdefgh|hhhhhhh
- BORDER_REFLECT: fedcba|abcdefgh|hgfedcb
- BORDER_REFLECT_101: gfedcb|abcdefgh|gfedcba
- BORDER_WRAP: cdefgh|abcdefgh|abcdefg
- BORDER_CONSTANT: 0000|abcdefgh|0000

In all cases except 1, an $\text{ksize} \times \text{ksize}$ separable kernel will be used to calculate the derivative. When $\text{ksize} = 1$, a 3×1 or 1×3 kernel will be used (i.e. no Gaussian smoothing is done). $\text{ksize} = 1$ can only be used for the first or the second x- or y- derivatives.

There is also the special value $\text{ksize} = \text{CV_SCHARR}$ (-1) that corresponds to a 3×3 Scharr filter that may give more accurate results than a 3×3 Sobel. The Scharr aperture is

$$\begin{bmatrix} -3 & 0 & 3 \\ -10 & 0 & 10 \\ -3 & 0 & 3 \end{bmatrix}$$

for the x-derivative or transposed for the y-derivative.

The function calculates the image derivative by convolving the image with the appropriate kernel:

$$\text{dst} = \frac{\partial^{xorder+yorder} \text{src}}{\partial x^{xorder} \partial y^{yorder}}$$

The Sobel operators combine Gaussian smoothing and differentiation, so the result is more or less resistant to the noise. Most often, the function is called with ($xorder = 1$, $yorder = 0$, $ksize = 3$) or ($xorder = 0$, $yorder = 1$, $ksize = 3$)

to calculate the first x- or y- image derivative. The first case corresponds to a kernel of:

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

and the second one corresponds to a kernel of:

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

2.3 2.3 Geometric Image Transformations

2.3.1 cv::ocl::resize

Resizes an image.

Supports CV_8UC1, CV_8UC4, CV_32FC1 and CV_32FC4 data types.

src Source image

dst Destination image. It will have size dsize (when it is non-zero) or the size computed from src.size() and fx and fy. The type of dst will be the same as of src.

dsize The destination image size. If it is zero, then it is computed as:

$\text{dsize} = \text{Size}(\text{round}(\text{fx} * \text{src.cols}), \text{round}(\text{fy} * \text{src.rows}))$

. Either dsize or both fx or fy must be non-zero.

fx The scale factor along the horizontal axis. When 0, it is computed as

$(\text{double})\text{dsize.width}/\text{src.cols}$

fy The scale factor along the vertical axis. When 0, it is computed as

$(\text{double})\text{dsize.height}/\text{src.rows}$

interpolation The interpolation method:

INTER_NEAREST nearest-neighbor interpolation

INTER_LINEAR bilinear interpolation (used by default)

INTER_AREA resampling using pixel area relation. It may be the preferred method for image decimation, as it gives moire-free results. But when the image is zoomed, it is similar to the INTER_NEAREST method

INTER_CUBIC bicubic interpolation over 4x4 pixel neighborhood

INTER_LANCZOS4 Lanczos interpolation over 8x8 pixel neighborhood

The function resize resizes an image src down to or up to the specified size. Note that the initial dst type or size are not taken into account. Instead the size and type are derived from the src, dsize, fx and fy. If you want to resize src so that it fits the pre-created dst, you may call the function as:

If you want to decimate the image by factor of 2 in each direction, you can call the function this way:

2.3.2 cv::ocl::warpAffine

Applies an affine transformation to an image.

Supports INTER_NEAREST, INTER_LINEAR, INTER_CUBIC types.

src Source image

dst Destination image; will have size dsize and the same type as src

M 2×3 transformation matrix

dsize Size of the destination image

flags A combination of interpolation methods, see cv::resize, and the optional flag WARP_INVERSE_MAP that means that M is the inverse transformation (dst→src)

The function warpAffine transforms the source image using the specified matrix:

$$\text{dst}(x, y) = \text{src}(M_{11}x + M_{12}y + M_{13}, M_{21}x + M_{22}y + M_{23})$$

when the flag WARP_INVERSE_MAP is set. Otherwise, the transformation is first inverted with cv::invertAffineTransform and then put in the formula above instead of M. The function can not operate in-place.

2.3.3 cv::ocl::warpPerspective

Applies a perspective transformation to an image.

Supports INTER_NEAREST, INTER_LINEAR, INTER_CUBIC types.

src Source image

dst Destination image; will have size dsize and the same type as src

M 3×3 transformation matrix

dsize Size of the destination image

flags A combination of interpolation methods, see cv::resize, and the optional flag WARP_INVERSE_MAP that means that M is the inverse transformation (dst→src)

The function warpPerspective transforms the source image using the specified matrix:

$$\text{dst}(x, y) = \text{src}\left(\frac{M_{11}x + M_{12}y + M_{13}}{M_{31}x + M_{32}y + M_{33}}, \frac{M_{21}x + M_{22}y + M_{23}}{M_{31}x + M_{32}y + M_{33}}\right)$$

when the flag WARP_INVERSE_MAP is set. Otherwise, the transformation is first inverted with cv::invert and then put in the formula above instead of M. The function can not operate in-place.

2.4 2.4 Miscellaneous Image Transformations

2.4.1 cv::ocl::cvtColor ()

Converts image from one color space to another.

Supports.

src The source image, 8-bit unsigned, 16-bit unsigned (CV_16UC...) or single-precision floatingpoint

dst The destination image; will have the same size and the same depth as src

code The color space conversion code; see the discussion

dstCn The number of channels in the destination image; if the parameter is 0, the number of the channels will be derived automatically from src and the code

The function converts the input image from one color space to another. In the case of transformation to/from RGB color space the ordering of the channels should be specified explicitly (RGB or BGR).

The conventional ranges for R, G and B channel values are:

- 0 to 255 for CV_8U images
- 0 to 65535 for CV_16U images and
- 0 to 1 for CV_32F images.

Of course, in the case of linear transformations the range does not matter, but in the non-linear cases the input RGB image should be normalized to the proper value range in order to get the correct results, e.g. for RGB→L*u*v* transformation. For example, if you have a 32-bit floatingpoint image directly converted from 8-bit image without any scaling, then it will have 0..255 value range, instead of the assumed by the function 0..1. So, before calling cvtColor, you need first to scale the image down:

The function can do the following transformations:

- Transformations within RGB space like adding/removing the alpha channel, reversing the channel order, conversion to/from 16-bit RGB color (R5:G6:B5 or R5:G5:B5), as well as

$$\text{RGB}[A] \text{ to Gray: } Y \leftarrow 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B$$

and

$$\text{Gray to RGB}[A] : R \leftarrow Y, G \leftarrow Y, B \leftarrow Y, A \leftarrow 0$$

The conversion from a RGB image to gray is done with:

- RGB←→CIE XYZ.Rec709 with D65 white point (CV_BGR2XYZ, CV_RGB2XYZ, CV_XYZ2BGR, CV_XYZ2RGB):

$$\begin{bmatrix} X \\ Y \\ Z \end{bmatrix} \leftarrow \begin{bmatrix} 0.412453 & 0.357580 & 0.180423 \\ 0.212671 & 0.715160 & 0.072169 \\ 0.019334 & 0.119193 & 0.950227 \end{bmatrix} \cdot \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

$$\begin{bmatrix} R \\ G \\ B \end{bmatrix} \leftarrow \begin{bmatrix} 3.240479 & -1.53715 & -0.498535 \\ -0.969256 & 1.875991 & 0.041556 \\ 0.055648 & -0.204043 & 1.057311 \end{bmatrix} \cdot \begin{bmatrix} X \\ Y \\ Z \end{bmatrix}$$

X, Y and Z cover the whole value range (in the case of floating-point images Z may exceed 1).

- RGB \leftrightarrow YCrCb JPEG (a.k.a. YCC) (CV_BGR2YCrCb, CV_RGB2YCrCb, CV_YCrCb2BGR, CV_YCrCb2RGB)

$$\begin{aligned} Y &\leftarrow 0.299 \cdot R + 0.587 \cdot G + 0.114 \cdot B \\ Cr &\leftarrow (R - Y) \cdot 0.713 + \text{delta} \\ Cb &\leftarrow (B - Y) \cdot 0.564 + \text{delta} \\ R &\leftarrow Y + 1.403 \cdot (Cr - \text{delta}) \\ G &\leftarrow Y - 0.344 \cdot (Cr - \text{delta}) - 0.714 \cdot (Cb - \text{delta}) \\ B &\leftarrow Y + 1.773 \cdot (Cb - \text{delta}) \end{aligned}$$

where

$$\text{delta} = \begin{cases} 128 & \text{for 8-bit images} \\ 32768 & \text{for 16-bit images} \\ 0.5 & \text{for floating-point images} \end{cases}$$

Y, Cr and Cb cover the whole value range.

- RGB \leftrightarrow HSV (CV_BGR2HSV, CV_RGB2HSV, CV_HSV2BGR, CV_HSV2RGB)
in the case of 8-bit and 16-bit images R, G and B are converted to floating-point format and scaled to fit the 0 to 1 range

$$\begin{aligned} Y &\leftarrow \max(R, G, B) \\ S &\leftarrow \begin{cases} \frac{V - \min(R, G, B)}{V} & \text{if } V \neq 0 \\ 0 & \text{otherwise} \end{cases} \\ H &\leftarrow \begin{cases} 60(G - B/S) & \text{if } V = R \\ 120 + 60(B - R)/S & \text{if } V = G \\ 240 + 60(R - G)/S & \text{if } V = B \end{cases} \end{aligned}$$

if $H < 0$ then $H \leftarrow H + 360$ On output $0 \leq L \leq 1$, $0 \leq S \leq 1$, $0 \leq H \leq 360$.

The values are then converted to the destination data type:

8-bit images

$V \leftarrow 255 \cdot V$, $S \leftarrow 255 \cdot S$, $H \leftarrow H/2$ (to fit to 0 to 255)

16-bit images (currently not supported)

$V < -65535 \cdot V$, $S < -65535 \cdot S$, $H < -H$

32-bit images H, S, V are left as is

- $RGB \leftarrow \rightarrow CIE\ L^*a^*b^*$ (CV_BGR2Lab, CV_RGB2Lab, CV_Lab2BGR, CV_Lab2RGB)
in the case of 8-bit and 16-bit images R, G and B are converted to floating-point format and scaled to fit the 0 to 1 range

$$\begin{bmatrix} X \\ Y \\ Z \end{bmatrix} \leftarrow \begin{bmatrix} 0.412453 & 0.357580 & 0.180423 \\ 0.212671 & 0.715160 & 0.072169 \\ 0.019334 & 0.119193 & 0.950227 \end{bmatrix} \cdot \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

$X \leftarrow X/X_n$, where $X_n=0.950456$
 $Z \leftarrow Z/Z_n$, where $Z_n=1.088754$

$$L \leftarrow \begin{cases} 116 * Y^{1/3} - 16 & \text{for } Y > 0.008856 \\ 903.3 * Y & \text{for } Y \leq 0.008856 \end{cases}$$

$$\begin{aligned} a &\leftarrow 500(f(X) - f(Y)) + \text{delta} \\ b &\leftarrow 200(f(Y) - f(Z)) + \text{delta} \end{aligned}$$

where

$$f(t) = \begin{cases} t^{1/3} & \text{for } t > 0.008856 \\ 7.787t + 16/116 & \text{for } t \leq 0.008856 \end{cases}$$

and

$$\text{delta} = \begin{cases} 128 & \text{for 8-bit images} \\ 0 & \text{for floating-point images} \end{cases}$$

On output $0 \leq L \leq 100$, $-127 \leq a \leq 127$, $-127 \leq b \leq 127$

The values are then converted to the destination data type:

8-bit images

$L \leftarrow L * 255 / 100$, $a \leftarrow a + 128$, $b \leftarrow b + 128$

16-bit images currently not supported

32-bit images L, a, b are left as is

- $RGB \leftarrow \rightarrow CIE\ L^*u^*v^*$ (CV_BGR2Luv, CV_RGB2Luv, CV_Luv2BGR, CV_Luv2RGB)
in the case of 8-bit and 16-bit images R, G and B are converted to floating-point format and scaled to fit 0 to 1 range

$$\begin{bmatrix} X \\ Y \\ Z \end{bmatrix} \leftarrow \begin{bmatrix} 0.412453 & 0.357580 & 0.180423 \\ 0.212671 & 0.715160 & 0.072169 \\ 0.019334 & 0.119193 & 0.950227 \end{bmatrix} \cdot \begin{bmatrix} R \\ G \\ B \end{bmatrix}$$

$$L \leftarrow \begin{cases} 116 * Y^{1/3} - 16 & \text{for } Y > 0.008856 \\ 903.3 * Y & \text{for } Y \leq 0.008856 \end{cases}$$

$$u' \leftarrow 4 * X / (X + 15 * Y + 3Z)$$

$$v' \leftarrow 9 * Y / (X + 15 * Y + 3Z)$$

$$u \leftarrow 13 * L * (u' - u_n) \quad \text{where } u_n = 0.19793943$$

$$v \leftarrow 13 * L * (v' - v_n) \quad \text{where } v_n = 0.46831096$$

On output $0 \leq L \leq 100$, $-134 \leq u \leq 220$, $-140 \leq v \leq 122$.

The values are then converted to the destination data type:

8-bit images

$L \leftarrow 255/100L$, $u \leftarrow 255/354(u + 134)$, $v \leftarrow 255/256(v + 140)$

16-bit images currently not supported

32-bit images L, u, v are left as is

The above formulas for converting RGB to/from various color spaces have been taken from multiple sources on Web, primarily from the Charles Poynton site <http://www.poynton.com/ColorFAQ.html>

- Bayer→RGB (CV_BayerBG2BGR, CV_BayerGB2BGR, CV_BayerRG2BGR, CV_BayerGR2BGR, CV_BayerBG2RGB, CV_BayerGB2RGB, CV_BayerRG2RGB, CV_BayerGR2RGB) The Bayer pattern is widely used in CCD and CMOS cameras. It allows one to get color pictures from a single plane where R,G and B pixels (sensors of a particular component) are interleaved like this:

<i>R</i>	<i>G</i>	<i>R</i>	<i>G</i>	<i>R</i>
<i>G</i>	<i>B</i>	<i>G</i>	<i>B</i>	<i>G</i>
<i>R</i>	<i>G</i>	<i>R</i>	<i>G</i>	<i>R</i>
<i>G</i>	<i>B</i>	<i>G</i>	<i>B</i>	<i>G</i>
<i>R</i>	<i>G</i>	<i>R</i>	<i>G</i>	<i>R</i>

The output RGB components of a pixel are interpolated from 1, 2 or 4 neighbors of the pixel having the same color. There are several modifications of the above pattern that can be achieved by shifting the pattern one pixel left and/or one pixel up. The two letters C1 and C2 in the conversion constants CV_Bayer C1C2 2BGR and CV_Bayer C1C2 2RGB indicate the particular pattern type - these are components from the second row, second and third columns, respectively. For example, the above pattern has very popular “BG” type.

2.4.2 cv::ocl::integral

Calculates the integral of an image.

Supports only CV_8UC1 source type.

image The source image, $W \times H$, 8-bit or floating-point (32f or 64f)

sum The integral image, $(W + 1) \times (H + 1)$, 32-bit integer or floating-point (32f or 64f)

sqsum The integral image for squared pixel values, $(W + 1) \times (H + 1)$, double precision floatingpoint (64f)

The functions integral calculate one or more integral images for the source image as following:

$$\text{sum}(X, Y) = \sum_{x < X, y < Y} \text{image}(x, y)$$

$$\text{sqsum}(X, Y) = \sum_{x < X, y < Y} \text{image}(x, y)^2$$

$$\text{tilted}(X, Y) = \sum_{y < Y, \text{abs}(x - X + 1) \leq Y - y - 1} \text{image}(x, y)$$

Using these integral images, one may calculate sum, mean and standard deviation over a specific up-right or rotated rectangular region of the image in a constant time, for example:

$$\sum_{x_1 \leq x < x_2, y_1 \leq y < y_2} \text{image}(x, y) = \text{sum}(x_2, y_2) - \text{sum}(x_1, y_2) - \text{sum}(x_2, y_1) + \text{sum}(x_1, y_1)$$

It makes possible to do a fast blurring or fast block correlation with variable window size, for example. In the case of multi-channel images, sums for each channel are accumulated independently.

As a practical example, the next figure shows the calculation of the integral of a straight rectangle Rect(3,3,3,2) and of a tilted rectangle Rect(5,1,2,3). The selected pixels in the original image are shown, as well as the relative pixels in the integral images sum and tilted.

$$P_0 = \{y, x\} = \{4, 4\}$$

$$P_1 = \{y, x + w\} = \{4, 7\}$$

$$P_2 = \{y + h, x\} = \{6, 4\}$$


$$P_3 = \{y + h, x + w\} = \{6, 7\}$$


$$P_0 = \{y, x\} = \{1, 5\}$$

$$P_1 = \{y + h, x + w\} = \{3, 7\}$$

$$P_2 = \{y + h, x - h\} = \{4, 2\}$$

$$P_3 = \{y + w + h, x + w - h\} = \{6, 4\}$$





2.4.3 cv::ocl::threshold

Applies a fixed-level threshold to each array element

Supports CV_8UC1 and CV_32FC1 data types.

src Source array (single-channel, 8-bit or 32-bit floating point)

dst Destination array; will have the same size and the same type as src

thresh Threshold value

maxVal Maximum value to use with THRESH_BINARY and THRESH_BINARY_INV thresholding types

thresholdType Thresholding type (see the discussion)

The function applies fixed-level thresholding to a single-channel array. The function is typically used to get a bi-level (binary) image out of a grayscale image or for removing a noise, i.e. filtering out pixels with too small or too large values. There are several types of thresholding that the function supports that are determined by thresholdType:

THRESH_BINARY

$$\text{dst}(x, y) = \begin{cases} \text{max Val} & \text{if } \text{src}(x, y) \geq \text{thresh} \\ 0 & \text{otherwise} \end{cases}$$

THRESH_BINARY_INV

$$\text{dst}(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) \geq \text{thresh} \\ \text{max Val} & \text{otherwise} \end{cases}$$

THRESH_TRUNC

$$\text{dst}(x, y) = \begin{cases} \text{threshold} & \text{if } \text{src}(x, y) \geq \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$$

THRESH_TOZERO

$$\text{dst}(x, y) = \begin{cases} \text{src}(x, y) & \text{if } \text{src}(x, y) \geq \text{thresh} \\ 0 & \text{otherwise} \end{cases}$$

THRESH TOZERO INV

$$\text{dst}(x, y) = \begin{cases} 0 & \text{if } \text{src}(x, y) \leq \text{thresh} \\ \text{src}(x, y) & \text{otherwise} \end{cases}$$

Also, the special value THRESH OTSU may be combined with one of the above values. In this case the function determines the optimal threshold value using Otsu's algorithm and uses it instead of the specified thresh. The function returns the computed threshold value. Currently, Otsu's method is implemented only for 8-bit images.

